Homework 5

Jiaxi li

Table of contents

Question 1	2
svm was the best so far Question 2	16
svm was the highest accuracy Question 3	29
as the number of hidden layers increases, the neural network becomes more expressive	50

! Important

Please read the instructions carefully before submitting your assignment.

- 1. This assignment requires you to only upload a PDF file on Canvas
- 2. Don't collapse any code cells before submitting.
- 3. Remember to make sure all your code output is rendered properly before uploading your submission.

Please add your name to the author information in the frontmatter before submitting your assignment

In this assignment, we will explore decision trees, support vector machines and neural networks for classification and regression. The assignment is designed to test your ability to fit and analyze these models with different configurations and compare their performance.

We will need the following packages:

```
packages <- c(
  "tibble",
  "dplyr",
  "readr",
  "tidyr",
  "purrr",
  "broom",
  "magrittr",
  "corrplot",
  "caret",
  "rpart",
  "rpart.plot",
  "e1071",
  "torch",
  "luz"
# renv::install(packages)
sapply(packages, require, character.only=T)
```

Question 1



Prediction of Median House prices

1.1 (2.5 points)

The data folder contains the housing.csv dataset which contains housing prices in California from the 1990 California census. The objective is to predict the median house price for California districts based on various features.

Read the data file as a tibble in R. Preprocess the data such that:

- 1. the variables are of the right data type, e.g., categorical variables are encoded as factors
- 2. all column names to lower case for consistency
- 3. Any observations with missing values are dropped

```
path <- "data/housing.csv"</pre>
  df <- read_csv(path) %>%
    mutate_if(is.character, as.factor) %>%
    rename_with(tolower) %>%
    drop_na()
Rows: 20640 Columns: 10
-- Column specification ------
Delimiter: ","
chr (1): ocean_proximity
dbl (9): longitude, latitude, housing_median_age, total_rooms, total_bedroom...
i Use `spec()` to retrieve the full column specification for this data.
i Specify the column types or set `show_col_types = FALSE` to quiet this message.
  df$ocean_proximity <- as.factor(df$ocean_proximity)</pre>
  head(df) # Insert your code here
# A tibble: 6 x 10
  longitude latitude housing_median_age total_rooms total_bedrooms population
      <dbl>
              <dbl>
                                 <dbl>
                                             <dbl>
                                                            <dbl>
                                                                      <dbl>
     -122.
                                    41
1
               37.9
                                              880
                                                             129
                                                                        322
2
     -122.
                                    21
                                              7099
                                                             1106
               37.9
                                                                       2401
3
     -122.
               37.8
                                    52
                                              1467
                                                             190
                                                                        496
4
     -122.
               37.8
                                    52
                                              1274
                                                             235
                                                                        558
     -122.
               37.8
                                    52
                                              1627
                                                             280
                                                                        565
     -122.
               37.8
                                    52
                                               919
                                                             213
                                                                        413
```

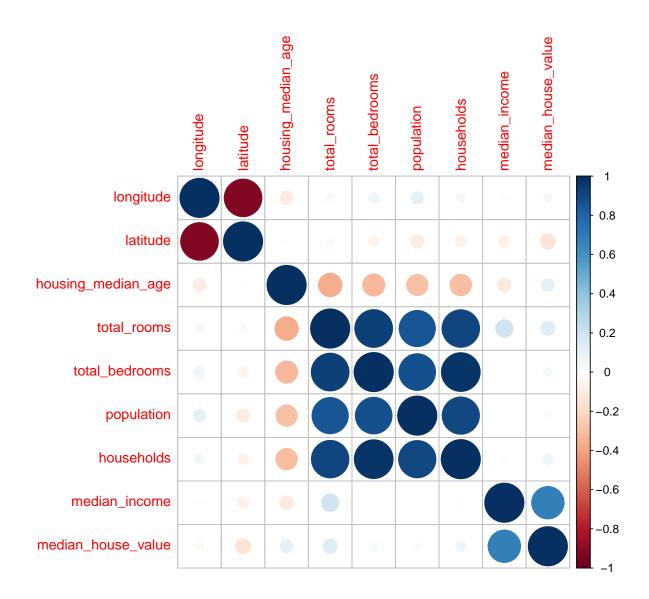
1.2 (2.5 points)

Visualize the correlation matrix of all numeric columns in df using corrplot()

```
numeric_columns <- df %>% select_if(is.numeric)
cor_matrix <- cor(numeric_columns)</pre>
```

[#] i 4 more variables: households <dbl>, median_income <dbl>,

[#] median_house_value <dbl>, ocean_proximity <fct>



1.3 (5 points)

Split the data df into df_train and df_split using test_ind in the code below:

```
set.seed(42)
test_ind <- sample(
   1:nrow(df),
   floor( nrow(df)/10 ),
   replace=FALSE
)

df_train <- df[-test_ind, ] # Insert your code here
df_test <- df[test_ind, ] # Insert your code here</pre>
```

1.4 (5 points)

Fit a linear regression model to predict the median_house_value:

- latitude
- longitude
- housing_median_age
- total_rooms
- total_bedrooms
- population
- median_income
- ocean_proximity

Interpret the coefficients and summarize your results.

```
(Intercept)
                         -2.273e+06 9.138e+04 -24.873 < 2e-16 ***
                         -2.681e+04 1.060e+03 -25.305 < 2e-16 ***
longitude
latitude
                         -2.539e+04 1.047e+03 -24.244 < 2e-16 ***
                          1.074e+03 4.616e+01 23.261 < 2e-16 ***
housing_median_age
total rooms
                         -6.159e+00 8.431e-01 -7.306 2.87e-13 ***
total_bedrooms
                          1.353e+02 4.254e+00 31.804 < 2e-16 ***
population
                         -3.413e+01 9.838e-01 -34.694 < 2e-16 ***
median_income
                          3.936e+04 3.573e+02 110.154 < 2e-16 ***
                         -4.018e+04 1.836e+03 -21.891 < 2e-16 ***
ocean_proximityINLAND
ocean_proximityISLAND
                          1.324e+05 3.442e+04
                                                3.847 0.00012 ***
                         -2.522e+03 2.022e+03 -1.247 0.21226
ocean_proximityNEAR BAY
ocean_proximityNEAR OCEAN 4.349e+03 1.658e+03
                                                2.622 0.00875 **
               0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
Signif. codes:
Residual standard error: 68780 on 18378 degrees of freedom
Multiple R-squared: 0.643, Adjusted R-squared: 0.6428
F-statistic: 3009 on 11 and 18378 DF, p-value: < 2.2e-16
```

1.5 (5 points)

Complete the rmse function for computing the Root Mean-Squared Error between the true y and the predicted yhat, and use it to compute the RMSE for the regression model on df_test

```
rmse <- function(y, yhat) {
   sqrt(mean((y - yhat)^2))
}

lm_predictions <- predict(lm_fit,newdata = df_test) # Insert your code here
rmse_lm <- rmse(df_test$median_house_value, lm_predictions)
print(rmse_lm)</pre>
```

[1] 68339.82

1.6 (5 points)

Fit a decision tree model to predict the median_house_value using the same predictors as in 1.4. Use the rpart() function.

```
rpart_fit <- rpart(median_house_value ~ . - households, data = df_train) # Insert your cod
  summary(rpart_fit) # Insert your code here
Call:
rpart(formula = median_house_value ~ . - households, data = df_train)
  n = 18390
          CP nsplit rel error
                                 xerror
                  0 1.0000000 1.0000592 0.011312769
1 0.30492555
                  1 0.6950744 0.6979744 0.008826537
2 0.12731459
                  2 0.5677599 0.5701815 0.007578389
3 0.05953042
                  3 0.5082294 0.5107395 0.007457761
4 0.04048333
                  4 0.4677461 0.4720499 0.007116065
5 0.01381152
                 5 0.4539346 0.4570311 0.007005307
6 0.01380725
7 0.01315679
                  6 0.4401273 0.4489693 0.006943697
8 0.01000000
                  7 0.4269706 0.4277617 0.006761393
Variable importance
     median income
                                                                  longitude
                      ocean proximity
                                                latitude
                                   22
housing_median_age
                          total rooms
                                              population
                 2
Node number 1: 18390 observations,
                                     complexity param=0.3049256
  mean=206394.2, MSE=1.324089e+10
  left son=2 (14519 obs) right son=3 (3871 obs)
  Primary splits:
                      < 5.03515 to the left, improve=0.30492560, (0 missing)
      median_income
      ocean_proximity splits as RLRRR,
                                               improve=0.23381360, (0 missing)
                                 to the right, improve=0.06666687, (0 missing)
      latitude
                      < 37.935
                      < -121.865 to the right, improve=0.03779756, (0 missing)
      longitude
                      < 2015.5 to the left, improve=0.02725354, (0 missing)
      total_rooms
  Surrogate splits:
                             to the left, agree=0.791, adj=0.005, (0 split)
      total rooms < 13799
                                      complexity param=0.1273146
Node number 2: 14519 observations,
  mean=173584.7, MSE=8.384653e+09
  left son=4 (5255 obs) right son=5 (9264 obs)
  Primary splits:
      ocean_proximity
                                                  improve=0.25465680, (0 missing)
                         splits as RLRRR,
                         < 3.12885 to the left, improve=0.16224220, (0 missing)
      median_income
      latitude
                         < 37.925
                                    to the right, improve=0.06160970, (0 missing)
```

```
to the left, improve=0.03829309, (0 missing)
      housing_median_age < 51.5
                         < -121.865 to the right, improve=0.03606499, (0 missing)
      longitude
 Surrogate splits:
                         < 37.955
                                    to the right, agree=0.727, adj=0.246, (0 split)
      latitude
     housing_median_age < 15.5
                                    to the left, agree=0.664, adj=0.071, (0 split)
                         < -116.905 to the right, agree=0.659, adj=0.057, (0 split)
      longitude
      population
                         < 343.5
                                    to the left, agree=0.649, adj=0.029, (0 split)
      median_income
                         < 1.79275 to the left, agree=0.644, adj=0.016, (0 split)
Node number 3: 3871 observations,
                                     complexity param=0.05953042
 mean=329452.8, MSE=1.22743e+10
 left son=6 (2701 obs) right son=7 (1170 obs)
 Primary splits:
                         < 6.81955 to the left, improve=0.30508300, (0 missing)
      median_income
                                                  improve=0.12921480, (0 missing)
      ocean_proximity
                         splits as
                                    RL-RR,
      housing_median_age < 27.5
                                    to the left, improve=0.08620557, (0 missing)
      latitude
                         < 37.965
                                    to the right, improve=0.06037955, (0 missing)
                         < -118.035 to the right, improve=0.03922245, (0 missing)
      longitude
  Surrogate splits:
      total bedrooms < 35
                                to the right, agree=0.702, adj=0.015, (0 split)
      population
                     < 64
                                to the right, agree=0.702, adj=0.014, (0 split)
                                to the right, agree=0.702, adj=0.013, (0 split)
      total rooms
                     < 208
Node number 4: 5255 observations,
                                     complexity param=0.01315679
 mean=112232.1, MSE=2.870763e+09
 left son=8 (3027 obs) right son=9 (2228 obs)
 Primary splits:
      median_income
                         < 3.03555 to the left,
                                                  improve=0.21236300, (0 missing)
                                    to the right, improve=0.04093847, (0 missing)
      latitude
                         < 34.825
      longitude
                         < -118.295 to the left, improve=0.03080534, (0 missing)
                         < 2296
                                    to the left, improve=0.02651350, (0 missing)
      total_rooms
      housing_median_age < 17.5
                                    to the right, improve=0.02167716, (0 missing)
  Surrogate splits:
                                    to the right, agree=0.631, adj=0.129, (0 split)
      housing_median_age < 14.5
                                    to the left, agree=0.623, adj=0.110, (0 split)
      total rooms
                         < 3012.5
                                    to the left, agree=0.597, adj=0.049, (0 split)
      population
                         < 2761.5
      total bedrooms
                         < 970
                                    to the left, agree=0.593, adj=0.041, (0 split)
      latitude
                         < 34.705
                                    to the right, agree=0.589, adj=0.030, (0 split)
Node number 5: 9264 observations,
                                     complexity param=0.04048333
 mean=208387, MSE=8.166e+09
  left son=10 (4061 obs) right son=11 (5203 obs)
 Primary splits:
```

```
improve=0.13030680, (0 missing)
      median_income
                         < 3.12285 to the left,
      longitude
                         < -118.305 to the right, improve=0.07512986, (0 missing)
                                    to the left, improve=0.03855985, (0 missing)
      housing_median_age < 51.5
                         < 33.995
                                    to the left,
                                                  improve=0.03798139, (0 missing)
      latitude
      total rooms
                         < 2032
                                    to the left, improve=0.02982308, (0 missing)
  Surrogate splits:
      total rooms
                         < 1280.5 to the left, agree=0.600, adj=0.088, (0 split)
      latitude
                         < 32.775
                                    to the left, agree=0.580, adj=0.041, (0 split)
                         < -123.09 to the left, agree=0.572, adj=0.024, (0 split)
      longitude
      housing_median_age < 37.5
                                   to the right, agree=0.572, adj=0.024, (0 split)
                                    to the left, agree=0.567, adj=0.012, (0 split)
      total_bedrooms
                         < 125.5
Node number 6: 2701 observations,
                                     complexity param=0.01381152
  mean=289177.6, MSE=8.861805e+09
  left son=12 (479 obs) right son=13 (2222 obs)
  Primary splits:
      ocean_proximity
                         splits as RL-RR,
                                                  improve=0.14050550, (0 missing)
      housing_median_age < 36.5
                                    to the left, improve=0.10106890, (0 missing)
      median income
                         < 5.78465 to the left, improve=0.07031948, (0 missing)
      latitude
                         < 37.955
                                    to the right, improve=0.05584316, (0 missing)
      longitude
                         < -122.105 to the right, improve=0.04390701, (0 missing)
  Surrogate splits:
                         < 38.335
      latitude
                                    to the right, agree=0.857, adj=0.192, (0 split)
                                    to the left, agree=0.827, adj=0.023, (0 split)
      housing_median_age < 2.5
      longitude
                         < -116.62 to the right, agree=0.826, adj=0.019, (0 split)
Node number 7: 1170 observations
  mean=422430.1, MSE=7.762749e+09
Node number 8: 3027 observations
  mean=91049.03, MSE=1.619675e+09
Node number 9: 2228 observations
  mean=141011.9, MSE=3.132597e+09
Node number 10: 4061 observations
  mean=171463.8, MSE=6.028421e+09
Node number 11: 5203 observations,
                                      complexity param=0.01380725
  mean=237205.9, MSE=7.939788e+09
  left son=22 (2182 obs) right son=23 (3021 obs)
  Primary splits:
      longitude
                         < -118.275 to the right, improve=0.08138477, (0 missing)
```

```
housing_median_age < 47.5
                                  to the left,
                                               improve=0.06867000, (0 missing)
   latitude
                       < 33.985
                                               improve=0.04674031, (0 missing)
                                  to the left,
                                               improve=0.02487397, (0 missing)
   median_income
                       < 4.13035 to the left,
   total_bedrooms
                       < 418.5
                                  to the left,
                                               improve=0.02369109, (0 missing)
Surrogate splits:
   latitude
                       < 34.145
                                  to the left,
                                               agree=0.884, adj=0.723, (0 split)
   ocean proximity
                       splits as L-RRR,
                                                agree=0.597, adj=0.040, (0 split)
                                  to the left, agree=0.591, adj=0.024, (0 split)
   housing_median_age < 7.5
   population
                       < 3681.5
                                  to the right, agree=0.582, adj=0.003, (0 split)
   total_rooms
                                  to the right, agree=0.582, adj=0.003, (0 split)
                       < 11924
```

Node number 12: 479 observations mean=213177.9, MSE=6.03936e+09

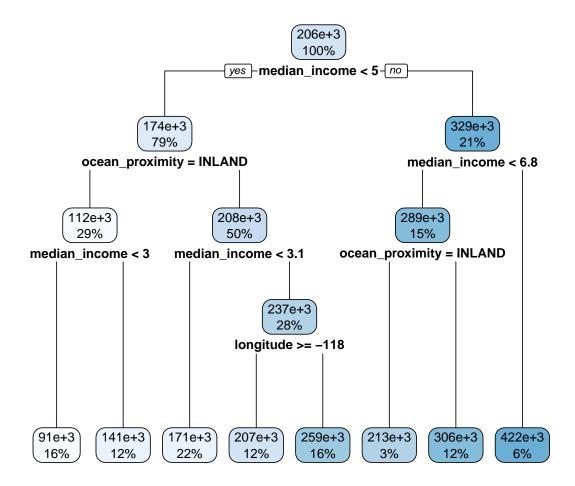
Node number 13: 2222 observations mean=305561, MSE=7.956696e+09

Node number 22: 2182 observations mean=207295.4, MSE=5.051886e+09

Node number 23: 3021 observations mean=258809.6, MSE=8.912757e+09

Visualize the decision tree using the rpart.plot() function.

```
rpart.plot(rpart_fit) # Insert your code here
```



Report the root mean squared error on the test set.

```
rpart_predictions <- predict(rpart_fit, newdata = df_test) # Insert your code here
rmse_dt <- rmse(df_test$median_house_value, rpart_predictions)
print(rmse_dt)</pre>
```

[1] 75876.87

1.7 (5 points)

Fit a support vector machine model to predict the median_house_value using the same predictors as in 1.4. Use the svm() function and use any kernel of your choice. Report the root mean squared error on the test set.

```
svm_fit <- svm(median_house_value ~ . - households, data = df_train) # Insert your code he
svm_predictions <- predict(svm_fit, newdata = df_test) # Insert your code here
svm_rmse <- rmse(df_test$median_house_value, svm_predictions)
print(svm_rmse)</pre>
```

[1] 56678.84

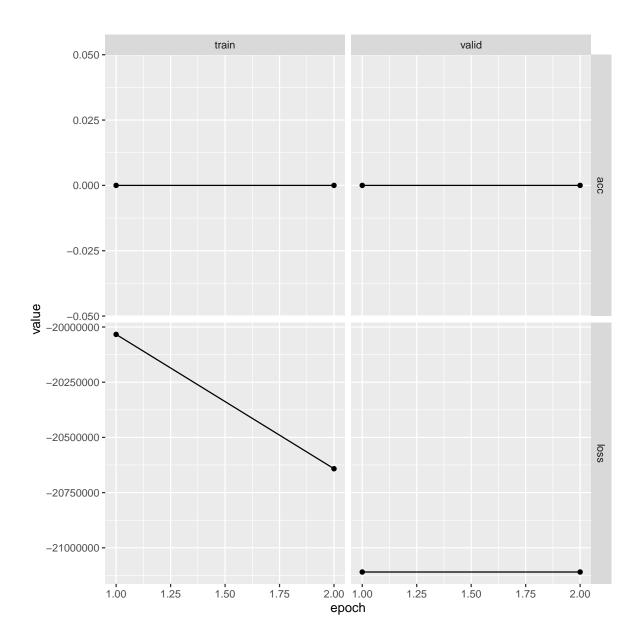
1.8 (25 points)

Initialize a neural network model architecture:

```
nn <- nn_module(</pre>
  initialize = function(p,q1,q2,q3) {
    self$hidden1 <- nn_linear(p, q1)</pre>
    self$hidden2 <- nn_linear(q1, q2)</pre>
    self$hidden3 <- nn_linear(q2, q3)</pre>
    self$output <- nn_linear(q3, 1) # Single output node for binary classification
    self$activation <- nn_relu()</pre>
    self$sigmoid <- nn_sigmoid()</pre>
  },
  forward = function(x) {
    x %>%
      self$hidden1() %>% self$activation() %>%
      self$hidden2() %>% self$activation() %>%
      self$hidden3() %>% self$activation() %>%
      self$output() %>% self$sigmoid() # Sigmoid activation function for binary classification
  }
)
```

Fit a neural network model to predict the median_house_value using the same predictors as in 1.4. Use the model.matrix function to create the covariate matrix and luz package for fitting the network with 32,16,8 nodes in each of the three hidden layers.

```
test_ind <- sample(1:nrow(df),23,replace = FALSE)</pre>
  M <- model.matrix(median_house_value ~ 0 + . , data = df_train)</pre>
  fit_nn <- nn %>%
    setup(loss = nn_bce_loss(),
           optimizer = optim_adam,
           metrics = list(luz_metric_accuracy())) %>%
    set_hparams(p = ncol(M),q1 = 32, q2 = 16, q3 = 8) %>%
    set_opt_hparams(lr= 0.001) %>%
    fit(data=list(
         model.matrix(median_house_value ~ 0 + .,data = df_train),df_train %>% select(median_
         valid_data = list(
         model.matrix(median_house_value ~ 0 + .,data = df_test),df_test %% select(median_house_value ~ 0 + .,data = df_test)
         epochs = 2,
         verbose = TRUE
         )
Epoch 1/2
Train metrics: Loss: -20033504 - Acc: 0
Valid metrics: Loss: -21109578 - Acc: 0
Epoch 2/2
Train metrics: Loss: -20641816 - Acc: 0
Valid metrics: Loss: -21109578 - Acc: 0
Plot the results of the training and validation loss and accuracy.
  plot(fit_nn) # Insert your code here
```



Report the root mean squared error on the test set.

```
test_result <- model.matrix(median_house_value ~ 0 + . , data = df_test)
nnet_predictions <- predict(fit_nn, test_result)</pre>
```

```
nnet_predictions <- as.array(nnet_predictions)</pre>
# Assuming `df_test$median_house_value` is your true values for the test set
rmse <- sqrt(mean((df_test$median_house_value - nnet_predictions)^2))</pre>
print(rmse)
```

[1] 242118.3



Warning

Remember to use the as_array() function to convert the predictions to a vector of numbers before computing the RMSE with rmse()

1.9 (5 points)

Summarize your results in a table comparing the RMSE for the different models. Which model performed best? Why do you think that is?

```
model_comparisons <- data.frame(</pre>
    Model = c("Linear Regression", "Decision Tree", "SVM", "nn_module"),
    RMSE = c(rmse_lm, rmse_dt, svm_rmse,rmse)
  )
  # Print the table
  model_comparisons
              Model
                         RMSE
1 Linear Regression 68339.82
2
      Decision Tree 75876.87
                SVM 56678.84
3
4
          nn_module 242118.32
```

knitr::kable(model_comparisons, caption = "Model Comparison based on RMSE")

Table 1: Model Comparison based on RMSE

Model	RMSE
Linear Regression	68339.82
Decision Tree	75876.87
SVM	56678.84
nn_module	242118.32

sym was the best so far

Question 2



9 50 points

Spam email classification

The data folder contains the spam.csv dataset. This dataset contains features extracted from a collection of spam and non-spam emails. The objective is to classify the emails as spam or non-spam.

2.1 (2.5 points)

Read the data file as a tibble in R. Preprocess the data such that:

- 1. the variables are of the right data type, e.g., categorical variables are encoded as factors
- 2. all column names to lower case for consistency
- 3. Any observations with missing values are dropped

```
spam_data <- read_csv("data/spambase.csv", show_col_types = FALSE) %>%
 mutate_if(is.character, as.factor) %>%
 rename_all(tolower) %>%
 drop_na()
```

2.2 (2.5 points)

Split the data df into df_train and df_split using test_ind in the code below:

```
set.seed(42)
test_ind <- sample(
   1:nrow(spam_data),
   floor( nrow(spam_data)/10 ),
   replace=FALSE
)

spam_train <- spam_data[-test_ind, ]
spam_test <- spam_data[test_ind, ]</pre>
```

Complete the overview function which returns a data frame with the following columns: accuracy, error, false positive rate, true positive rate, between the true true_class and the predicted pred_class for any classification model.

```
overview <- function(pred_class, true_class) {</pre>
 accuracy <- mean(pred_class == true_class)</pre>
 error <- mean(pred_class != true_class)</pre>
 true_positives <- sum(pred_class == 1 & true_class == 1)</pre>
 true_negatives <- sum(pred_class == 0 & true_class == 0)</pre>
 false_positives <- sum(pred_class == 1 & true_class == 0)</pre>
 false negatives <- sum(pred class == 0 & true class == 1)
 true_positive_rate <- true_positives / (true_positives + false_negatives)</pre>
 false positive rate <- false positives / (false positives + true negatives)
 return(
    data.frame(
      accuracy = accuracy,
      error = error,
      true_positive_rate = true_positive_rate,
      false_positive_rate = false_positive_rate
    )
 )
}
```

2.3 (5 points)

Fit a logistic regression model to predict the spam variable using the remaining predictors. Report the prediction accuracy on the test set.

```
glm_model <- glm(spam ~ ., family = binomial, data = spam_train)

Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

predictions <- predict(glm_model, newdata = spam_test, type = "response")
predicted_classes <- ifelse(predictions > 0.5, 1, 0)

accuracy1 <- mean(predicted_classes == spam_test)
print(accuracy1)

[1] 0.5123313</pre>
```

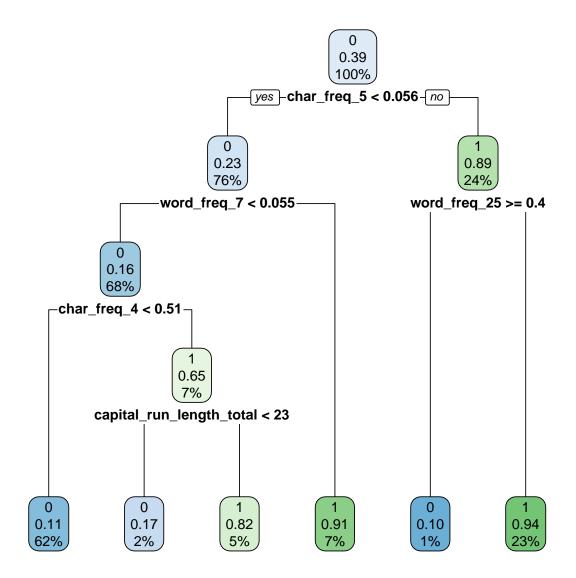
2.4 (5 points)

Fit a decision tree model to predict the spam variable using the remaining predictors. Use the rpart() function and set the method argument to "class".

```
rpart_classes <- rpart(spam ~ ., data = spam_train, method = "class") # Insert your code h</pre>
```

Visualize the decision tree using the rpart.plot() function.

```
rpart.plot(rpart_classes) # Insert your code here
```



Report the prediction accuracy on the test set.

```
rpart_predictions <- predict(rpart_classes, newdata = spam_test, type = "class")
rpart_predictions_num <- as.numeric(levels(rpart_predictions))[rpart_predictions]
rpart_accuracy <- mean(rpart_predictions_num == spam_test) # Insert your code here
print(rpart_accuracy)</pre>
```

[1] 0.5263118

2.5 (5 points)

Fit a support vector machine model to predict the spam variable using the remaining predictors. Use the sym() function and use any kernel of your choice. Remember to set the type argument to "C-classification" if you haven't already converted spam to be of type factor.

```
svm_fit <- svm(spam ~ ., data = spam_train, type = "C-classification", kernel = "radial")</pre>
```

Report the prediction accuracy on the test set.

```
svm_classes <- predict(svm_fit, newdata = spam_test) # Insert your code here</pre>
svm_accuracy <- mean(svm_classes == spam_test$spam)</pre>
# Print the accuracy
print(svm_accuracy)
```

[1] 0.923913

2.6 (25 points)

Using the same neural network architecture as in 1.9, fit a neural network model to predict the spam variable using the remaining predictors.

A Classification vs. Regression

Note that the neural network in Q 1.9 was a regression model. You will need to modify the neural network architecture to be a classification model by changing the output layer to have a single node with a sigmoid activation function.

Use the model.matrix function to create the covariate matrix and luz package for fitting the network with 32, 16, 8 nodes in each of the three hidden layers.

```
M2 <- model.matrix(spam ~ 0 + . , data = spam_train)</pre>
fit_nn2 <- nn %>%
  setup(loss = nn_bce_loss(),
        optimizer = optim_adam,
```

```
metrics = list(luz_metric_accuracy())) %>%
    set_hparams(p = ncol(M2),q1 = 32, q2 = 16, q3 = 8) %>%
    set_opt_hparams(lr= 0.001) %>%
    fit(data=list(
        model.matrix(spam ~ 0 + .,data = spam_train),spam_train %>% select(spam) %>% as.matr
        valid_data = list(
        model.matrix(spam ~ 0 + .,data = spam_test),spam_test %>% select(spam) %>% as.matrix
        epochs = 100,
        verbose = TRUE
        )
Epoch 1/100
Train metrics: Loss: 0.7532 - Acc: 12.5659
Valid metrics: Loss: 0.6081 - Acc: 12.8609
Epoch 2/100
Train metrics: Loss: 0.4521 - Acc: 12.5581
Valid metrics: Loss: 0.5522 - Acc: 12.8609
Epoch 3/100
Train metrics: Loss: 0.3537 - Acc: 12.5426
Valid metrics: Loss: 0.3075 - Acc: 12.8609
Epoch 4/100
Train metrics: Loss: 0.3727 - Acc: 12.5814
Valid metrics: Loss: 0.6486 - Acc: 12.8609
Epoch 5/100
Train metrics: Loss: 0.3093 - Acc: 12.5581
Valid metrics: Loss: 0.451 - Acc: 12.8609
Epoch 6/100
Train metrics: Loss: 0.2851 - Acc: 12.5581
Valid metrics: Loss: 0.4946 - Acc: 12.8609
Epoch 7/100
Train metrics: Loss: 0.232 - Acc: 12.5504
Valid metrics: Loss: 0.4382 - Acc: 12.8609
Epoch 8/100
Train metrics: Loss: 0.2709 - Acc: 12.5504
Valid metrics: Loss: 0.2461 - Acc: 12.8609
Epoch 9/100
Train metrics: Loss: 0.3106 - Acc: 12.5736
Valid metrics: Loss: 0.4305 - Acc: 12.8609
Epoch 10/100
Train metrics: Loss: 0.2446 - Acc: 12.5271
Valid metrics: Loss: 0.4824 - Acc: 12.8609
Epoch 11/100
```

```
Train metrics: Loss: 0.2616 - Acc: 12.5581
Valid metrics: Loss: 0.2398 - Acc: 12.8609
Epoch 12/100
Train metrics: Loss: 0.1983 - Acc: 12.5659
Valid metrics: Loss: 0.4451 - Acc: 12.8609
Epoch 13/100
Train metrics: Loss: 0.184 - Acc: 12.5659
Valid metrics: Loss: 0.4156 - Acc: 12.8609
Epoch 14/100
Train metrics: Loss: 0.2089 - Acc: 12.5659
Valid metrics: Loss: 0.2306 - Acc: 12.8609
Epoch 15/100
Train metrics: Loss: 0.2202 - Acc: 12.5504
Valid metrics: Loss: 0.4522 - Acc: 12.8609
Epoch 16/100
Train metrics: Loss: 0.1782 - Acc: 12.5426
Valid metrics: Loss: 0.4237 - Acc: 12.8609
Epoch 17/100
Train metrics: Loss: 0.2366 - Acc: 12.5814
Valid metrics: Loss: 0.4312 - Acc: 12.8609
Epoch 18/100
Train metrics: Loss: 0.1756 - Acc: 12.5504
Valid metrics: Loss: 0.4034 - Acc: 12.8609
Epoch 19/100
Train metrics: Loss: 0.2083 - Acc: 12.5426
Valid metrics: Loss: 0.276 - Acc: 12.8609
Epoch 20/100
Train metrics: Loss: 0.202 - Acc: 12.5581
Valid metrics: Loss: 0.4093 - Acc: 12.8609
Epoch 21/100
Train metrics: Loss: 0.2151 - Acc: 12.5504
Valid metrics: Loss: 0.224 - Acc: 12.8609
Epoch 22/100
Train metrics: Loss: 0.1671 - Acc: 12.5426
Valid metrics: Loss: 0.2155 - Acc: 12.8609
Epoch 23/100
Train metrics: Loss: 0.1904 - Acc: 12.5504
Valid metrics: Loss: 0.4119 - Acc: 12.8609
Epoch 24/100
Train metrics: Loss: 0.1747 - Acc: 12.5581
Valid metrics: Loss: 0.2215 - Acc: 12.8609
Epoch 25/100
Train metrics: Loss: 0.2382 - Acc: 12.5581
```

```
Valid metrics: Loss: 0.4598 - Acc: 12.8609
Epoch 26/100
Train metrics: Loss: 0.2182 - Acc: 12.5581
Valid metrics: Loss: 0.4114 - Acc: 12.8609
Epoch 27/100
Train metrics: Loss: 0.1738 - Acc: 12.5736
Valid metrics: Loss: 0.4961 - Acc: 12.8609
Epoch 28/100
Train metrics: Loss: 0.1847 - Acc: 12.5426
Valid metrics: Loss: 0.4022 - Acc: 12.8609
Epoch 29/100
Train metrics: Loss: 0.1566 - Acc: 12.5581
Valid metrics: Loss: 0.3983 - Acc: 12.8609
Epoch 30/100
Train metrics: Loss: 0.185 - Acc: 12.5581
Valid metrics: Loss: 0.4091 - Acc: 12.8609
Epoch 31/100
Train metrics: Loss: 0.1724 - Acc: 12.5426
Valid metrics: Loss: 0.4295 - Acc: 12.8609
Epoch 32/100
Train metrics: Loss: 0.1956 - Acc: 12.5814
Valid metrics: Loss: 0.4236 - Acc: 12.8609
Epoch 33/100
Train metrics: Loss: 0.1572 - Acc: 12.5659
Valid metrics: Loss: 0.406 - Acc: 12.8609
Epoch 34/100
Train metrics: Loss: 0.1624 - Acc: 12.5504
Valid metrics: Loss: 0.4036 - Acc: 12.8609
Epoch 35/100
Train metrics: Loss: 0.1515 - Acc: 12.5736
Valid metrics: Loss: 0.403 - Acc: 12.8609
Epoch 36/100
Train metrics: Loss: 0.1632 - Acc: 12.5659
Valid metrics: Loss: 0.5226 - Acc: 12.8609
Epoch 37/100
Train metrics: Loss: 0.1603 - Acc: 12.5581
Valid metrics: Loss: 0.3986 - Acc: 12.8609
Epoch 38/100
Train metrics: Loss: 0.1484 - Acc: 12.5426
Valid metrics: Loss: 0.4024 - Acc: 12.8609
Epoch 39/100
Train metrics: Loss: 0.1554 - Acc: 12.5504
Valid metrics: Loss: 0.4059 - Acc: 12.8609
```

```
Epoch 40/100
Train metrics: Loss: 0.1816 - Acc: 12.5426
Valid metrics: Loss: 0.4087 - Acc: 12.8609
Epoch 41/100
Train metrics: Loss: 0.1615 - Acc: 12.5426
Valid metrics: Loss: 0.472 - Acc: 12.8609
Epoch 42/100
Train metrics: Loss: 0.1519 - Acc: 12.5581
Valid metrics: Loss: 0.4347 - Acc: 12.8609
Epoch 43/100
Train metrics: Loss: 0.139 - Acc: 12.5736
Valid metrics: Loss: 0.4263 - Acc: 12.8609
Epoch 44/100
Train metrics: Loss: 0.1916 - Acc: 12.5659
Valid metrics: Loss: 0.2802 - Acc: 12.8609
Epoch 45/100
Train metrics: Loss: 0.1503 - Acc: 12.5736
Valid metrics: Loss: 0.4141 - Acc: 12.8609
Epoch 46/100
Train metrics: Loss: 0.1914 - Acc: 12.5581
Valid metrics: Loss: 0.4424 - Acc: 12.8609
Epoch 47/100
Train metrics: Loss: 0.1502 - Acc: 12.5426
Valid metrics: Loss: 0.4071 - Acc: 12.8609
Epoch 48/100
Train metrics: Loss: 0.1848 - Acc: 12.5426
Valid metrics: Loss: 0.4209 - Acc: 12.8609
Epoch 49/100
Train metrics: Loss: 0.1622 - Acc: 12.5891
Valid metrics: Loss: 0.4166 - Acc: 12.8609
Epoch 50/100
Train metrics: Loss: 0.1524 - Acc: 12.5581
Valid metrics: Loss: 0.4095 - Acc: 12.8609
Epoch 51/100
Train metrics: Loss: 0.1403 - Acc: 12.5659
Valid metrics: Loss: 0.4071 - Acc: 12.8609
Epoch 52/100
Train metrics: Loss: 0.1498 - Acc: 12.5736
Valid metrics: Loss: 0.4025 - Acc: 12.8609
Epoch 53/100
Train metrics: Loss: 0.1388 - Acc: 12.5581
Valid metrics: Loss: 0.4025 - Acc: 12.8609
Epoch 54/100
```

```
Train metrics: Loss: 0.1691 - Acc: 12.5581
Valid metrics: Loss: 0.2318 - Acc: 12.8609
Epoch 55/100
Train metrics: Loss: 0.1311 - Acc: 12.5659
Valid metrics: Loss: 0.3971 - Acc: 12.8609
Epoch 56/100
Train metrics: Loss: 0.1346 - Acc: 12.5736
Valid metrics: Loss: 0.397 - Acc: 12.8609
Epoch 57/100
Train metrics: Loss: 0.1406 - Acc: 12.5659
Valid metrics: Loss: 0.2118 - Acc: 12.8609
Epoch 58/100
Train metrics: Loss: 0.1251 - Acc: 12.5504
Valid metrics: Loss: 0.2247 - Acc: 12.8609
Epoch 59/100
Train metrics: Loss: 0.1355 - Acc: 12.5659
Valid metrics: Loss: 0.3976 - Acc: 12.8609
Epoch 60/100
Train metrics: Loss: 0.1427 - Acc: 12.5581
Valid metrics: Loss: 0.2122 - Acc: 12.8609
Epoch 61/100
Train metrics: Loss: 0.1258 - Acc: 12.5581
Valid metrics: Loss: 0.2142 - Acc: 12.8609
Epoch 62/100
Train metrics: Loss: 0.1292 - Acc: 12.5504
Valid metrics: Loss: 0.4086 - Acc: 12.8609
Epoch 63/100
Train metrics: Loss: 0.1214 - Acc: 12.5581
Valid metrics: Loss: 0.2201 - Acc: 12.8609
Epoch 64/100
Train metrics: Loss: 0.1267 - Acc: 12.5736
Valid metrics: Loss: 0.2206 - Acc: 12.8609
Epoch 65/100
Train metrics: Loss: 0.1274 - Acc: 12.5504
Valid metrics: Loss: 0.5663 - Acc: 12.8609
Epoch 66/100
Train metrics: Loss: 0.1293 - Acc: 12.5504
Valid metrics: Loss: 0.2279 - Acc: 12.8609
Epoch 67/100
Train metrics: Loss: 0.1204 - Acc: 12.5581
Valid metrics: Loss: 0.2345 - Acc: 12.8609
Epoch 68/100
Train metrics: Loss: 0.1235 - Acc: 12.5581
```

```
Valid metrics: Loss: 0.2234 - Acc: 12.8609
Epoch 69/100
Train metrics: Loss: 0.1288 - Acc: 12.5581
Valid metrics: Loss: 0.2455 - Acc: 12.8609
Epoch 70/100
Train metrics: Loss: 0.1218 - Acc: 12.5736
Valid metrics: Loss: 0.2335 - Acc: 12.8609
Epoch 71/100
Train metrics: Loss: 0.1247 - Acc: 12.5736
Valid metrics: Loss: 0.419 - Acc: 12.8609
Epoch 72/100
Train metrics: Loss: 0.119 - Acc: 12.5659
Valid metrics: Loss: 0.246 - Acc: 12.8609
Epoch 73/100
Train metrics: Loss: 0.1714 - Acc: 12.5349
Valid metrics: Loss: 0.2336 - Acc: 12.8609
Epoch 74/100
Train metrics: Loss: 0.1235 - Acc: 12.5814
Valid metrics: Loss: 0.4178 - Acc: 12.8609
Epoch 75/100
Train metrics: Loss: 0.1204 - Acc: 12.5736
Valid metrics: Loss: 0.3476 - Acc: 12.8609
Epoch 76/100
Train metrics: Loss: 0.1586 - Acc: 12.5504
Valid metrics: Loss: 0.4054 - Acc: 12.8609
Epoch 77/100
Train metrics: Loss: 0.1419 - Acc: 12.5581
Valid metrics: Loss: 0.4254 - Acc: 12.8609
Epoch 78/100
Train metrics: Loss: 0.1172 - Acc: 12.5659
Valid metrics: Loss: 0.4131 - Acc: 12.8609
Epoch 79/100
Train metrics: Loss: 0.1179 - Acc: 12.5659
Valid metrics: Loss: 0.4218 - Acc: 12.8609
Epoch 80/100
Train metrics: Loss: 0.1119 - Acc: 12.5349
Valid metrics: Loss: 0.4189 - Acc: 12.8609
Epoch 81/100
Train metrics: Loss: 0.1079 - Acc: 12.5659
Valid metrics: Loss: 0.4195 - Acc: 12.8609
Epoch 82/100
Train metrics: Loss: 0.1209 - Acc: 12.5659
Valid metrics: Loss: 0.2373 - Acc: 12.8609
```

```
Epoch 83/100
Train metrics: Loss: 0.1122 - Acc: 12.5504
Valid metrics: Loss: 0.407 - Acc: 12.8609
Epoch 84/100
Train metrics: Loss: 0.1365 - Acc: 12.5504
Valid metrics: Loss: 0.4046 - Acc: 12.8609
Epoch 85/100
Train metrics: Loss: 0.1346 - Acc: 12.5504
Valid metrics: Loss: 0.2421 - Acc: 12.8609
Epoch 86/100
Train metrics: Loss: 0.1061 - Acc: 12.5736
Valid metrics: Loss: 0.2282 - Acc: 12.8609
Epoch 87/100
Train metrics: Loss: 0.1167 - Acc: 12.5349
Valid metrics: Loss: 0.2423 - Acc: 12.8609
Epoch 88/100
Train metrics: Loss: 0.1038 - Acc: 12.5504
Valid metrics: Loss: 0.2251 - Acc: 12.8609
Epoch 89/100
Train metrics: Loss: 0.1072 - Acc: 12.5349
Valid metrics: Loss: 0.2406 - Acc: 12.8609
Epoch 90/100
Train metrics: Loss: 0.0999 - Acc: 12.5581
Valid metrics: Loss: 0.23 - Acc: 12.8609
Epoch 91/100
Train metrics: Loss: 0.1066 - Acc: 12.5426
Valid metrics: Loss: 0.2506 - Acc: 12.8609
Epoch 92/100
Train metrics: Loss: 0.1215 - Acc: 12.5504
Valid metrics: Loss: 0.421 - Acc: 12.8609
Epoch 93/100
Train metrics: Loss: 0.1053 - Acc: 12.5504
Valid metrics: Loss: 0.2723 - Acc: 12.8609
Epoch 94/100
Train metrics: Loss: 0.097 - Acc: 12.5426
Valid metrics: Loss: 0.2511 - Acc: 12.8609
Epoch 95/100
Train metrics: Loss: 0.1205 - Acc: 12.5349
Valid metrics: Loss: 0.2265 - Acc: 12.8609
Epoch 96/100
Train metrics: Loss: 0.0977 - Acc: 12.5736
Valid metrics: Loss: 0.238 - Acc: 12.8609
Epoch 97/100
```

```
Train metrics: Loss: 0.1005 - Acc: 12.5349
Valid metrics: Loss: 0.2579 - Acc: 12.8609
Epoch 98/100
Train metrics: Loss: 0.0949 - Acc: 12.5814
Valid metrics: Loss: 0.2371 - Acc: 12.8609
Epoch 99/100
Train metrics: Loss: 0.0958 - Acc: 12.5426
Valid metrics: Loss: 0.2775 - Acc: 12.8609
Epoch 100/100
Train metrics: Loss: 0.0933 - Acc: 12.5504
Valid metrics: Loss: 0.2335 - Acc: 12.8609
  spam_test_result <- model.matrix(spam ~ 0 + . , data = spam_test)</pre>
  nnet2_predictions <- predict(fit_nn2, spam_test_result)</pre>
  nnet_predictions <- as.array(nnet2_predictions)</pre>
  # Assuming `df_test$median_house_value` is your true values for the test set
  nn_accuracy <- mean(nnet_predictions == spam_test$spam)</pre>
  print(nn_accuracy)
[1] 0.08043478
```

2.7 (5 points)

Summarize your results in a table comparing the accuracy metrics for the different models.

```
model_comparisons2 <- data.frame(
    Model = c("logistic regression model", "decision tree model", "SVM","nn"),
    RMSE = c(accuracy1, rpart_accuracy, svm_accuracy,nn_accuracy)
)

# Print the table
print(model_comparisons2)</pre>
```

Model RMSE

```
1 logistic regression model 0.51233133
2
        decision tree model 0.52631184
3
                        SVM 0.92391304
4
                         nn 0.08043478
```

```
knitr::kable(model_comparisons2, caption = "Model Comparison based on RMSE")
```

Table 2: Model Comparison based on RMSE

Model	RMSE
logistic regression model	0.5123313
decision tree model	0.5263118
SVM	0.9239130
nn	0.0804348
decision tree model SVM	0.5263118 0.9239130

svm was the highest accuracy

If you were to choose a model to classify spam emails, which model would you choose? Think about the context of the problem and the cost of false positives and false negatives.

Question 3



9 60 points

Three spirals classification

To better illustrate the power of depth in neural networks, we will use a toy dataset called the "Three Spirals" data. This dataset consists of two intertwined spirals, making it challenging for shallow models to classify the data accurately.



This is a multi-class classification problem

The dataset can be generated using the provided R code below:

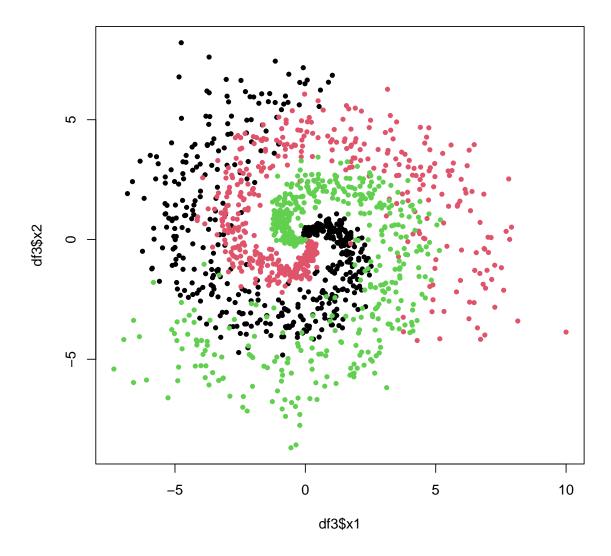
```
generate_three_spirals <- function(){</pre>
 set.seed(42)
 n <- 500
 noise <- 0.2
 t <- (1:n) / n * 2 * pi
  x1 <- c(
     t * (sin(t) + rnorm(n, 0, noise)),
      t * (sin(t + 2 * pi/3) + rnorm(n, 0, noise)),
      t * (sin(t + 4 * pi/3) + rnorm(n, 0, noise))
 x2 <- c(
     t * (cos(t) + rnorm(n, 0, noise)),
      t * (cos(t + 2 * pi/3) + rnorm(n, 0, noise)),
     t * (cos(t + 4 * pi/3) + rnorm(n, 0, noise))
 y <- as.factor(
   c(
     rep(0, n),
     rep(1, n),
     rep(2, n)
  )
 return(tibble::tibble(x1=x1, x2=x2, y=y))
```

3.1 (5 points)

Generate the three spirals dataset using the code above. Plot x_1 vs x_2 and use the y variable to color the points.

```
df3 <- generate_three_spirals()

plot(
   df3$x1, df3$x2,
   col = df3$y,
   pch = 20
)</pre>
```



Define a grid of 100 points from -10 to 10 in both x_1 and x_2 using the expand.grid(). Save it as a tibble called df_test.

```
x1 <- seq(-10, 10, length.out = 100)
x2 <- seq(-10, 10, length.out = 100)
grid <- expand.grid(x1 = x1, x2 = x2)
df_test3 <- as_tibble(grid)</pre>
```

3.2 (10 points)

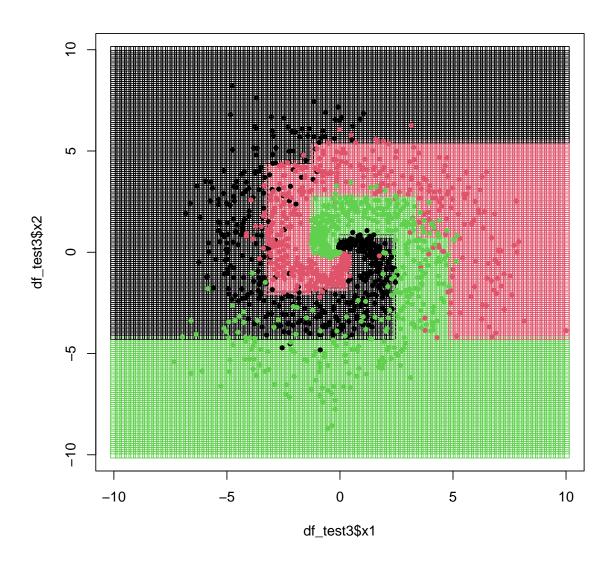
Fit a classification tree model to predict the y variable using the x1 and x2 predictors, and plot the decision boundary.

```
rpart_fit <- rpart(y ~ x1 + x2, data = df3, method = "class")
rpart_classes <- predict(rpart_fit, newdata = df_test3, type = "class")</pre>
```

Plot the decision boundary using the following function:

```
plot_decision_boundary <- function(predictions){
  plot(
    df_test3$x1, df_test3$x2,
    col = predictions,
    pch = 0
)
  points(
    df3$x1, df3$x2,
    col = df3$y,
    pch = 20
)
}</pre>
```

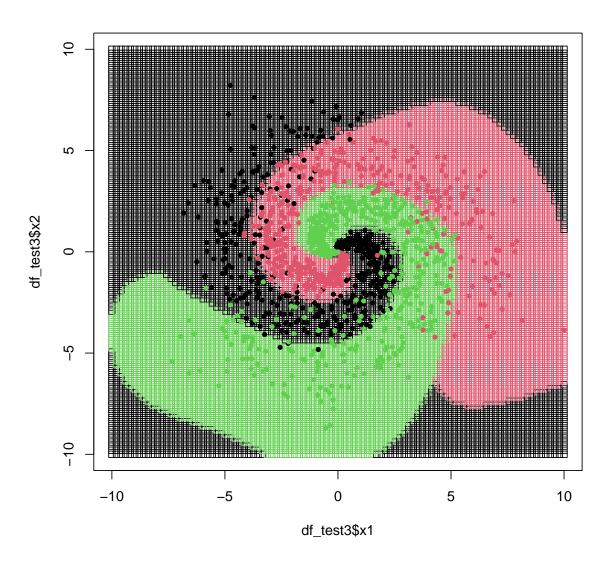
plot_decision_boundary(rpart_classes)



3.3 (10 points)

Fit a support vector machine model to predict the y variable using the x1 and x2 predictors. Use the svm() function and use any kernel of your choice. Remember to set the type argument to "C-classification" if you haven't converted y to be of type factor.

```
df3$y <- as.factor(df3$y)
svm_fit <- svm(y ~ x1 + x2, data = df3, type = "C-classification", kernel = "radial") # In
svm_classes <- predict(svm_fit, newdata = df_test3) # Insert your code here
plot_decision_boundary(svm_classes)</pre>
```



⚠ Instructions

For the next questions, you will need to fit a series of neural networks. In all cases, you can:

- set the number of units in each hidden layer to 10
- set the output dimension o to 3 (remember this is multinomial classification)
- use the appropriate loss function for the problem (not nn_bce_loss)
- set the number of epochs to 50
- fit the model using the luz package

You can use any optimizer of your choice, but you will need to tune the learning rate for each problem.

3.4 (10 points)

Fit a neural network with 1 hidden layer to predict the y variable using the x1 and x2 predictors.

```
NN1 <- nn_module(</pre>
  initialize = function(p, q1, o) {
    self$hidden1 <- nn_linear(p, q1)</pre>
    self$output <- nn_linear(q1, o)</pre>
    self$activation <- nn_relu()</pre>
  },
  forward = function(x) {
    x %>%
      self$hidden1() %>%
      self$activation() %>%
      self$output()
  }
)
fit_1 <- NN1 %>%
  setup(
    loss = nn_cross_entropy_loss(),
    optimizer = optim_adam,
    metrics = list(luz_metric_accuracy())
  ) %>%
```

```
set_hparams(
      p = 2, # Number of input features (x1 and x2)
      q1 = 10, # Number of hidden units
      o = 3 # Number of output classes
    ) %>%
    set_opt_hparams(lr= 0.005) %>%
    fit(
      data = list(
        x = df3 \%% select(x1, x2) %>% as.matrix(),
        y = df3$y %>% as.integer()
      ),
      epochs = 50,
      verbose = TRUE
    )
Epoch 1/50
Train metrics: Loss: 1.0672 - Acc: 0.3906
Epoch 2/50
Train metrics: Loss: 0.9795 - Acc: 0.3641
Epoch 3/50
Train metrics: Loss: 0.9367 - Acc: 0.4076
Epoch 4/50
Train metrics: Loss: 0.8973 - Acc: 0.4735
Epoch 5/50
Train metrics: Loss: 0.8558 - Acc: 0.5543
Epoch 6/50
Train metrics: Loss: 0.8221 - Acc: 0.5951
Epoch 7/50
Train metrics: Loss: 0.791 - Acc: 0.5965
Epoch 8/50
Train metrics: Loss: 0.7616 - Acc: 0.6488
Epoch 9/50
Train metrics: Loss: 0.7345 - Acc: 0.6311
Epoch 10/50
Train metrics: Loss: 0.7089 - Acc: 0.6556
Epoch 11/50
Train metrics: Loss: 0.6833 - Acc: 0.6732
Epoch 12/50
Train metrics: Loss: 0.6606 - Acc: 0.6936
Epoch 13/50
Train metrics: Loss: 0.6459 - Acc: 0.6902
Epoch 14/50
```

Train metrics: Loss: 0.6309 - Acc: 0.7113

Epoch 15/50

Train metrics: Loss: 0.6154 - Acc: 0.7079

Epoch 16/50

Train metrics: Loss: 0.6014 - Acc: 0.7167

Epoch 17/50

Train metrics: Loss: 0.5916 - Acc: 0.7181

Epoch 18/50

Train metrics: Loss: 0.5838 - Acc: 0.7133

Epoch 19/50

Train metrics: Loss: 0.5738 - Acc: 0.7385

Epoch 20/50

Train metrics: Loss: 0.5658 - Acc: 0.7262

Epoch 21/50

Train metrics: Loss: 0.5569 - Acc: 0.7405

Epoch 22/50

Train metrics: Loss: 0.5532 - Acc: 0.7371

Epoch 23/50

Train metrics: Loss: 0.546 - Acc: 0.7534

Epoch 24/50

Train metrics: Loss: 0.5395 - Acc: 0.7473

Epoch 25/50

Train metrics: Loss: 0.5387 - Acc: 0.7357

Epoch 26/50

Train metrics: Loss: 0.5336 - Acc: 0.7541

Epoch 27/50

Train metrics: Loss: 0.5309 - Acc: 0.7534

Epoch 28/50

Train metrics: Loss: 0.5256 - Acc: 0.7554

Epoch 29/50

Train metrics: Loss: 0.5241 - Acc: 0.75

Epoch 30/50

Train metrics: Loss: 0.5141 - Acc: 0.7561

Epoch 31/50

Train metrics: Loss: 0.5156 - Acc: 0.7493

Epoch 32/50

Train metrics: Loss: 0.5161 - Acc: 0.7527

Epoch 33/50

Train metrics: Loss: 0.5103 - Acc: 0.7473

Epoch 34/50

Train metrics: Loss: 0.5094 - Acc: 0.7534

Epoch 35/50

Train metrics: Loss: 0.5079 - Acc: 0.7615

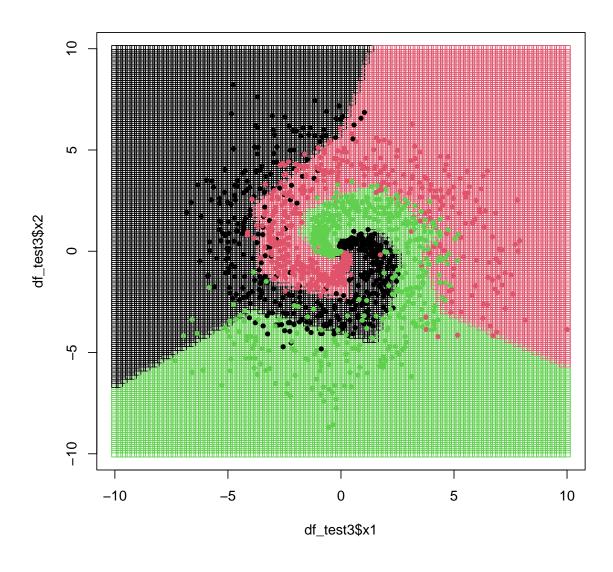
```
Epoch 36/50
Train metrics: Loss: 0.5016 - Acc: 0.7697
Epoch 37/50
Train metrics: Loss: 0.4982 - Acc: 0.7731
Epoch 38/50
Train metrics: Loss: 0.5033 - Acc: 0.7609
Epoch 39/50
Train metrics: Loss: 0.5062 - Acc: 0.7609
Epoch 40/50
Train metrics: Loss: 0.4949 - Acc: 0.7704
Epoch 41/50
Train metrics: Loss: 0.5013 - Acc: 0.7697
Epoch 42/50
Train metrics: Loss: 0.4894 - Acc: 0.7602
Epoch 43/50
Train metrics: Loss: 0.4979 - Acc: 0.7595
Epoch 44/50
Train metrics: Loss: 0.4908 - Acc: 0.7636
Epoch 45/50
Train metrics: Loss: 0.4882 - Acc: 0.7806
Epoch 46/50
Train metrics: Loss: 0.4919 - Acc: 0.7656
Epoch 47/50
Train metrics: Loss: 0.487 - Acc: 0.7785
Epoch 48/50
Train metrics: Loss: 0.4851 - Acc: 0.7595
Epoch 49/50
Train metrics: Loss: 0.4838 - Acc: 0.769
Epoch 50/50
Train metrics: Loss: 0.4882 - Acc: 0.769
```

In order to generate the class predictions, you will need to use the predict() function as follows

```
test_matrix <- df_test3 %>% select(x1, x2) %>% as.matrix
fit_1_predictions <- max.col(predict(fit_1, test_matrix))</pre>
```

Plot the results using the plot_decision_boundary() function.

```
plot_decision_boundary(fit_1_predictions)
```



3.5 (10 points)

Fit a neural network with 0 hidden layers to predict the y variable using the x1 and x2 predictors.

```
NNO <- nn_module(
    initialize = function(p, o) {
      self$output <- nn_linear(p, o)</pre>
      self$activation <- nn_relu()</pre>
    },
    forward = function(x) {
      x %>%
        self$activation() %>%
        self$output()
    }
  )
  setup(
      loss = nn_cross_entropy_loss(),
      optimizer = optim_adam,
      metrics = list(luz_metric_accuracy())
    ) %>%
    set_hparams(
      p = 2
      o = 3
    ) %>%
    set_opt_hparams(lr= 0.005) %>%
    fit(
      data = list(
        x = df3 \%% select(x1, x2) %% as.matrix(),
        y = df3$y %>% as.integer()
      ),
      epochs = 50,
      verbose = TRUE
    )
Epoch 1/50
Train metrics: Loss: 1.233 - Acc: 0.2996
Epoch 2/50
Train metrics: Loss: 1.1301 - Acc: 0.2908
Epoch 3/50
Train metrics: Loss: 1.0813 - Acc: 0.3166
Epoch 4/50
Train metrics: Loss: 1.065 - Acc: 0.3519
Epoch 5/50
Train metrics: Loss: 1.0578 - Acc: 0.3702
```

Epoch 6/50

Train metrics: Loss: 1.0553 - Acc: 0.3743

Epoch 7/50

Train metrics: Loss: 1.054 - Acc: 0.3757

Epoch 8/50

Train metrics: Loss: 1.0538 - Acc: 0.3838

Epoch 9/50

Train metrics: Loss: 1.0532 - Acc: 0.3784

Epoch 10/50

Train metrics: Loss: 1.0521 - Acc: 0.3947

Epoch 11/50

Train metrics: Loss: 1.0536 - Acc: 0.3818

Epoch 12/50

Train metrics: Loss: 1.0529 - Acc: 0.3601

Epoch 13/50

Train metrics: Loss: 1.055 - Acc: 0.3791

Epoch 14/50

Train metrics: Loss: 1.053 - Acc: 0.3764

Epoch 15/50

Train metrics: Loss: 1.0519 - Acc: 0.3811

Epoch 16/50

Train metrics: Loss: 1.0508 - Acc: 0.3879

Epoch 17/50

Train metrics: Loss: 1.0531 - Acc: 0.3723

Epoch 18/50

Train metrics: Loss: 1.0532 - Acc: 0.3716

Epoch 19/50

Train metrics: Loss: 1.0538 - Acc: 0.3689

Epoch 20/50

Train metrics: Loss: 1.0525 - Acc: 0.3764

Epoch 21/50

Train metrics: Loss: 1.0533 - Acc: 0.3689

Epoch 22/50

Train metrics: Loss: 1.0537 - Acc: 0.373

Epoch 23/50

Train metrics: Loss: 1.054 - Acc: 0.3784

Epoch 24/50

Train metrics: Loss: 1.0534 - Acc: 0.3662

Epoch 25/50

Train metrics: Loss: 1.0539 - Acc: 0.3832

Epoch 26/50

Train metrics: Loss: 1.052 - Acc: 0.3702

Epoch 27/50

```
Train metrics: Loss: 1.0533 - Acc: 0.377 Epoch 28/50
```

Train metrics: Loss: 1.0523 - Acc: 0.3655

Epoch 29/50

Train metrics: Loss: 1.0527 - Acc: 0.3852

Epoch 30/50

Train metrics: Loss: 1.0529 - Acc: 0.3798

Epoch 31/50

Train metrics: Loss: 1.0539 - Acc: 0.3757

Epoch 32/50

Train metrics: Loss: 1.0523 - Acc: 0.3696

Epoch 33/50

Train metrics: Loss: 1.0552 - Acc: 0.3655

Epoch 34/50

Train metrics: Loss: 1.0516 - Acc: 0.3872

Epoch 35/50

Train metrics: Loss: 1.0545 - Acc: 0.3832

Epoch 36/50

Train metrics: Loss: 1.0539 - Acc: 0.3635

Epoch 37/50

Train metrics: Loss: 1.0539 - Acc: 0.3716

Epoch 38/50

Train metrics: Loss: 1.0542 - Acc: 0.3709

Epoch 39/50

Train metrics: Loss: 1.0525 - Acc: 0.3743

Epoch 40/50

Train metrics: Loss: 1.0513 - Acc: 0.3743

Epoch 41/50

Train metrics: Loss: 1.0528 - Acc: 0.3764

Epoch 42/50

Train metrics: Loss: 1.0531 - Acc: 0.3798

Epoch 43/50

Train metrics: Loss: 1.0531 - Acc: 0.377

Epoch 44/50

Train metrics: Loss: 1.0542 - Acc: 0.3777

Epoch 45/50

Train metrics: Loss: 1.0541 - Acc: 0.3736

Epoch 46/50

Train metrics: Loss: 1.0522 - Acc: 0.3736

Epoch 47/50

Train metrics: Loss: 1.0524 - Acc: 0.3757

Epoch 48/50

Train metrics: Loss: 1.0531 - Acc: 0.3709

Epoch 49/50

Train metrics: Loss: 1.0517 - Acc: 0.3716

Epoch 50/50

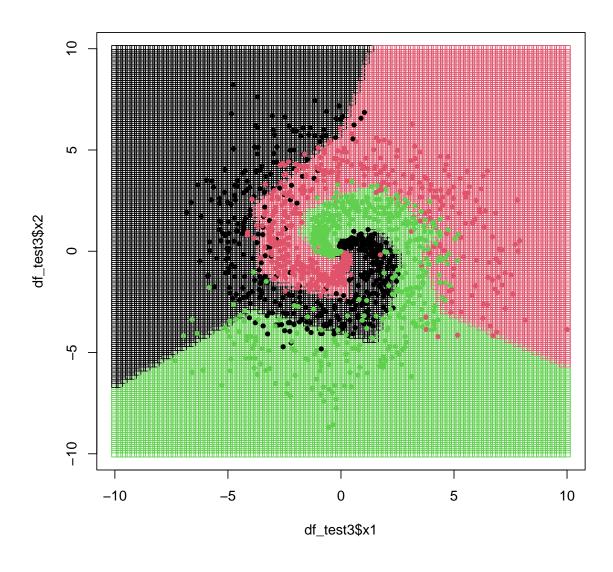
Train metrics: Loss: 1.0538 - Acc: 0.3859

Plot the results using the plot_decision_boundary() function.

```
fit_0_predictions <- max.col(predict(fit_0, test_matrix))</pre>
```

Plot the results using the plot_decision_boundary() function.

```
plot_decision_boundary(fit_1_predictions)
```



3.6 (10 points)

Fit a neural network with 3 hidden layers to predict the y variable using the x1 and x2 predictors.

```
NN2 <- nn_module(
  initialize = function(p, q1, q2, q3, o){
    self$hidden1 <- nn_linear(p, q1)</pre>
    self$hidden2 <- nn_linear(q1, q2)</pre>
    self$hidden3 <- nn_linear(q2, q3)</pre>
    self$output <- nn_linear(q3, o) # Single output node for binary classification</pre>
    self$activation <- nn_relu()</pre>
    self$sigmoid <- nn_sigmoid()</pre>
  },
  forward = function(x) {
    x %>%
      self$hidden1() %>% self$activation() %>%
      self$hidden2() %>% self$activation() %>%
      self$hidden3() %>% self$activation() %>%
      self$output() %>% self$sigmoid() # Sigmoid activation function for binary classification
  }
)
fit_2 <- NN2 %>%
  setup(
    loss = nn_cross_entropy_loss(),
    optimizer = optim_adam,
    metrics = list(luz_metric_accuracy())
  ) %>%
  set_hparams(
    p = 2,
    q1 = 10,
    q2 = 10,
    q3 = 10,
    o = 3
  ) %>%
  set_opt_hparams(lr= 0.005) %>%
  fit(
    data = list(
      x = df3 \%\% select(x1, x2) %>% as.matrix(),
      y = df3$y %>% as.integer()
    ),
    epochs = 50,
    verbose = TRUE
  )
```

Epoch 1/50

Train metrics: Loss: 1.0833 - Acc: 0.3899

Epoch 2/50

Train metrics: Loss: 1.0396 - Acc: 0.3995

Epoch 3/50

Train metrics: Loss: 0.9916 - Acc: 0.4769

Epoch 4/50

Train metrics: Loss: 0.9463 - Acc: 0.5312

Epoch 5/50

Train metrics: Loss: 0.8878 - Acc: 0.6603

Epoch 6/50

Train metrics: Loss: 0.8492 - Acc: 0.6821

Epoch 7/50

Train metrics: Loss: 0.8041 - Acc: 0.7086

Epoch 8/50

Train metrics: Loss: 0.7812 - Acc: 0.712

Epoch 9/50

Train metrics: Loss: 0.769 - Acc: 0.6895

Epoch 10/50

Train metrics: Loss: 0.7563 - Acc: 0.68

Epoch 11/50

Train metrics: Loss: 0.7509 - Acc: 0.6882

Epoch 12/50

Train metrics: Loss: 0.7429 - Acc: 0.6936

Epoch 13/50

Train metrics: Loss: 0.7355 - Acc: 0.7052

Epoch 14/50

Train metrics: Loss: 0.7358 - Acc: 0.697

Epoch 15/50

Train metrics: Loss: 0.7362 - Acc: 0.7065

Epoch 16/50

Train metrics: Loss: 0.7282 - Acc: 0.7126

Epoch 17/50

Train metrics: Loss: 0.7295 - Acc: 0.7181

Epoch 18/50

Train metrics: Loss: 0.7287 - Acc: 0.7174

Epoch 19/50

Train metrics: Loss: 0.7225 - Acc: 0.7201

Epoch 20/50

Train metrics: Loss: 0.719 - Acc: 0.7269

Epoch 21/50

Train metrics: Loss: 0.721 - Acc: 0.7255

Epoch 22/50

Train metrics: Loss: 0.7154 - Acc: 0.7276

Epoch 23/50

Train metrics: Loss: 0.7134 - Acc: 0.7364

Epoch 24/50

Train metrics: Loss: 0.7164 - Acc: 0.7317

Epoch 25/50

Train metrics: Loss: 0.7163 - Acc: 0.7371

Epoch 26/50

Train metrics: Loss: 0.715 - Acc: 0.7385

Epoch 27/50

Train metrics: Loss: 0.7109 - Acc: 0.7439

Epoch 28/50

Train metrics: Loss: 0.7106 - Acc: 0.7446

Epoch 29/50

Train metrics: Loss: 0.7135 - Acc: 0.7466

Epoch 30/50

Train metrics: Loss: 0.7152 - Acc: 0.7459

Epoch 31/50

Train metrics: Loss: 0.7072 - Acc: 0.7595

Epoch 32/50

Train metrics: Loss: 0.71 - Acc: 0.7486

Epoch 33/50

Train metrics: Loss: 0.7106 - Acc: 0.7548

Epoch 34/50

Train metrics: Loss: 0.7025 - Acc: 0.7643

Epoch 35/50

Train metrics: Loss: 0.7031 - Acc: 0.7704

Epoch 36/50

Train metrics: Loss: 0.702 - Acc: 0.7636

Epoch 37/50

Train metrics: Loss: 0.7032 - Acc: 0.7629

Epoch 38/50

Train metrics: Loss: 0.7041 - Acc: 0.7711

Epoch 39/50

Train metrics: Loss: 0.7 - Acc: 0.7724

Epoch 40/50

Train metrics: Loss: 0.6998 - Acc: 0.7765

Epoch 41/50

Train metrics: Loss: 0.7002 - Acc: 0.7704

Epoch 42/50

Train metrics: Loss: 0.6937 - Acc: 0.784

Epoch 43/50

Train metrics: Loss: 0.6997 - Acc: 0.7826

Epoch 44/50

```
Train metrics: Loss: 0.6967 - Acc: 0.7819
```

Epoch 45/50

Train metrics: Loss: 0.6909 - Acc: 0.7948

Epoch 46/50

Train metrics: Loss: 0.6986 - Acc: 0.7867

Epoch 47/50

Train metrics: Loss: 0.6858 - Acc: 0.8071

Epoch 48/50

Train metrics: Loss: 0.681 - Acc: 0.805

Epoch 49/50

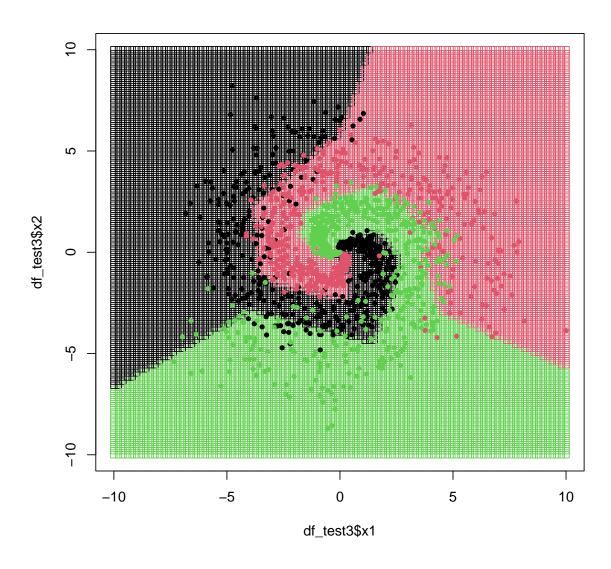
Train metrics: Loss: 0.6924 - Acc: 0.7894

Epoch 50/50

Train metrics: Loss: 0.6886 - Acc: 0.7969

Plot the results using the plot_decision_boundary() function.

```
fit_2_predictions <- max.col(predict(fit_2, test_matrix))
plot_decision_boundary(fit_1_predictions)</pre>
```



3.7 (5 points)

What are the differences between the models? How do the decision boundaries change as the number of hidden layers increases?

as the number of hidden layers increases, the neural network becomes more expressive

i Session Information

Print your R session information using the following command

sessionInfo()

R version 4.3.2 (2023-10-31 ucrt)

Platform: x86_64-w64-mingw32/x64 (64-bit)
Running under: Windows 11 x64 (build 22631)

Matrix products: default

locale:

- [1] LC_COLLATE=English_United States.utf8
- [2] LC_CTYPE=English_United States.utf8
- [3] LC_MONETARY=English_United States.utf8
- [4] LC_NUMERIC=C
- [5] LC_TIME=English_United States.utf8

time zone: America/New_York
tzcode source: internal

attached base packages:

[1] stats graphics grDevices utils datasets methods base

other attached packages:

[1]	luz_0.4.0	torch_0.12.0	e1071_1.7-14	rpart.plot_3.1.2
[5]	rpart_4.1.21	caret_6.0-94	lattice_0.21-9	ggplot2_3.4.4
[9]	corrplot_0.92	magrittr_2.0.3	broom_1.0.5	purrr_1.0.2
[13]	tidyr_1.3.1	readr_2.1.5	dplyr_1.1.4	tibble_3.2.1

farver_2.1.1

loaded via a namespace (and not attached):

[1] tidyselect_1.2.0 timeDate_4032.109

[4]	fastmap_1.1.1	pROC_1.18.5	digest_0.6.34
[7]	timechange_0.3.0	lifecycle_1.0.4	ellipsis_0.3.2
[10]	survival_3.5-7	processx_3.8.3	compiler_4.3.2
[13]	progress_1.2.3	rlang_1.1.3	tools_4.3.2
[16]	utf8_1.2.4	yam1_2.3.8	data.table_1.15.0
[19]	knitr_1.45	labeling_0.4.3	<pre>prettyunits_1.2.0</pre>

[22] bit_4.0.5	plyr_1.8.9	withr_3.0.0
[25] nnet_7.3-19	grid_4.3.2	$stats4_4.3.2$
[28] fansi_1.0.6	colorspace_2.1-0	future_1.33.1
[31] globals_0.16.3	scales_1.3.0	iterators_1.0.14
[34] MASS_7.3-60	zeallot_0.1.0	cli_3.6.2
[37] crayon_1.5.2	rmarkdown_2.25	generics_0.1.3
[40] rstudioapi_0.15.0	<pre>future.apply_1.11.1</pre>	reshape2_1.4.4
[43] tzdb_0.4.0	proxy_0.4-27	stringr_1.5.1
[46] splines_4.3.2	parallel_4.3.2	coro_1.0.4
[49] vctrs_0.6.5	hardhat_1.3.1	Matrix_1.6-1.1
[52] jsonlite_1.8.8	callr_3.7.3	hms_1.1.3
[55] bit64_4.0.5	listenv_0.9.1	foreach_1.5.2
[58] gower_1.0.1	recipes_1.0.10	glue_1.7.0
[61] parallelly_1.37.1	ps_1.7.6	codetools_0.2-19
[64] lubridate_1.9.3	stringi_1.8.3	gtable_0.3.4
[67] munsell_0.5.0	pillar_1.9.0	htmltools_0.5.7
[70] ipred_0.9-14	lava_1.8.0	R6_2.5.1
[73] vroom_1.6.5	evaluate_0.23	backports_1.4.1
[76] class_7.3-22	Rcpp_1.0.12	nlme_3.1-163
[79] prodlim_2023.08.28	xfun_0.41	fs_1.6.3
[82] pkgconfig_2.0.3	ModelMetrics_1.2.2.2	
	_	